

```
#dependency import
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
import numpy as np
```

## Data Import

```
#import data
data_full = pd.read_csv("./Pre-
Super_Day_candidate_dataset__28candidate_29 2.csv")
```

```
data_full.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	User ID	100000 non-null	object
1	applications	100000 non-null	int64
2	Reason	100000 non-null	object
3	Loan_Amount	100000 non-null	int64
4	FICO_score	100000 non-null	int64
5	Fico_Score_group	100000 non-null	object
6	Employment_Status	100000 non-null	object
7	Employment_Sector	93593 non-null	object
8	Monthly_Gross_Income	100000 non-null	int64
9	Monthly_Housing_Payment	100000 non-null	int64
10	Ever_Bankrupt_or_Foreclose	100000 non-null	int64
11	Lender	100000 non-null	object
12	Approved	100000 non-null	int64
13	bounty	100000 non-null	int64

```
dtypes: int64(8), object(6)
```

```
memory usage: 10.7+ MB
```

```
#data description
```

```
data_full.describe()
```

	applications	Loan_Amount	FICO_score	Monthly_Gross_Income
count	100000.0	100000.000000	100000.000000	100000.000000
mean	1.0	45234.350000	629.34961	5871.899350
std	0.0	28705.453665	88.66160	2882.939639
min	1.0	5000.000000	300.000000	2000.000000

25%	1.0	20000.000000	572.00000	3704.000000
50%	1.0	40000.000000	634.00000	5172.500000
75%	1.0	70000.000000	693.00000	7631.000000
max	1.0	100000.000000	850.00000	19997.000000

	Monthly_Housing_Payment	Ever_Bankrupt_or_Foreclose
Approved \		
count	100000.000000	100000.000000
mean	1649.693970	0.022460
std	623.443127	0.148175
min	300.000000	0.000000
25%	1229.000000	0.000000
50%	1665.000000	0.000000
75%	2046.000000	0.000000
max	3300.000000	1.000000

	bounty
count	100000.000000
mean	26.415000
std	78.385644
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	350.000000

## Data preprocessing

- Dropping Features: I have decided to drop `User_ID` columns because it is unique for every data point. I also drop the `applications` column because it is redundant.
- Label Encoding: I have also decided to label encode `object` (i.e. `String`) columns, and keep track of which categories these encodings fall into.

```
#drop useless features
to_drop = ["User ID", "applications", "Reason", "Employment_Sector"]
data= data_full.drop (data_full [to_drop], axis=1)
```

```

#encode non-numerical features

from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()

categorical_columns =
["Fico_Score_group", "Employment_Status", "Lender"]

#transform each column
for column in categorical_columns:
    data[column] = label_encoder.fit_transform(data[column])

```

## 1. Explore the variables relationship with approvability:


Feature relevance was assessed using correlation analysis and Random Forest feature importance.

```

#define correlation matrix.
correlation_matrix = data.drop("bounty", axis = 1).corr()

#visualize correlation
plt.figure(figsize = (16, 9))
sns.heatmap(correlation_matrix, annot = True, vmin = -1, vmax = 1)
plt.show()

```

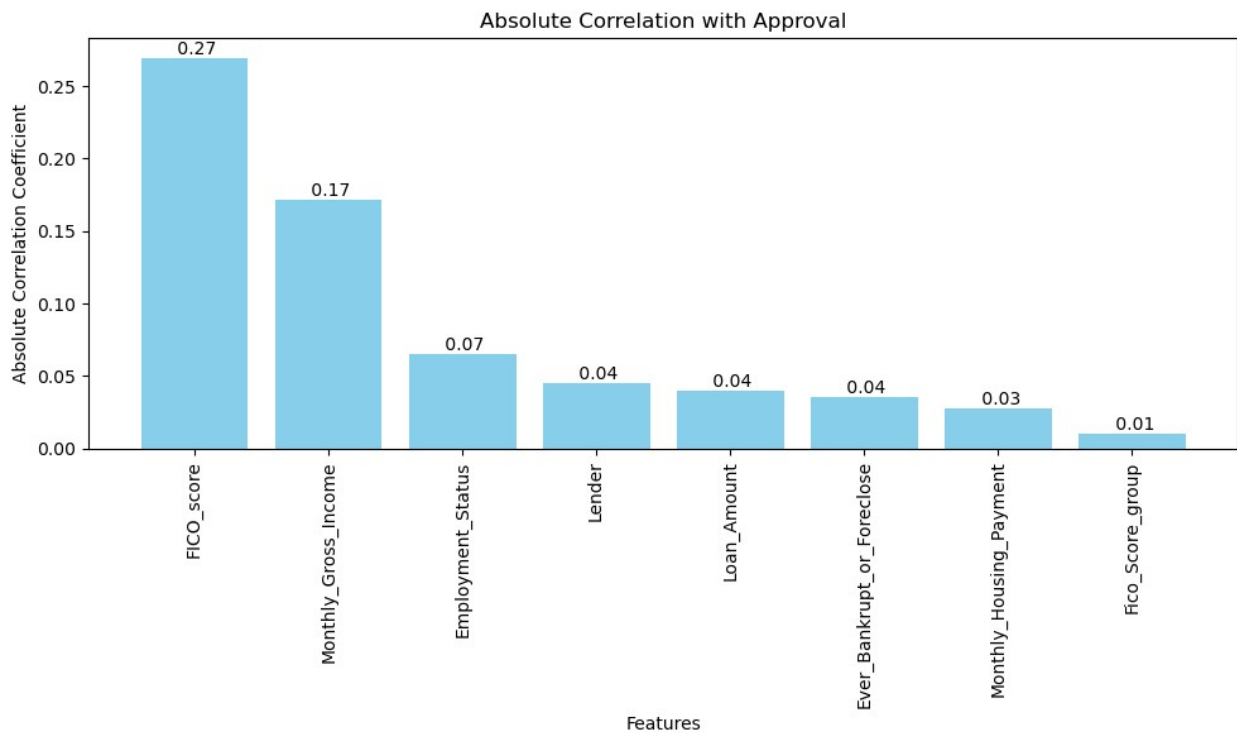


```
# rank correlation with approval based on absolute value of
correlation_with_approved =
correlation_matrix["Approved"].abs().sort_values(ascending = False)

correlation_with_approved = correlation_with_approved.drop("Approved")
#plot results
plt.figure(figsize=(10,6))
bars =
plt.bar(correlation_with_approved.index,correlation_with_approved.valu
es,color = "skyblue" )
plt.title("Absolute Correlation with Approval")
plt.xlabel("Features")
plt.ylabel("Absolute Correlation Coefficient")
plt.xticks(rotation = 90)
plt.tight_layout()

# Annotate bars with their corresponding coefficents for better
readability.
for bar in bars :
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width() / 2, height,
```

```
f'{height: .2f}', ha = 'center', va = 'bottom')
plt.show()
```



### Random Forest Feature Importance

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split

#dataset split
X = data.drop(["Approved", "bounty"], axis = 1)
y = data["Approved"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 1)

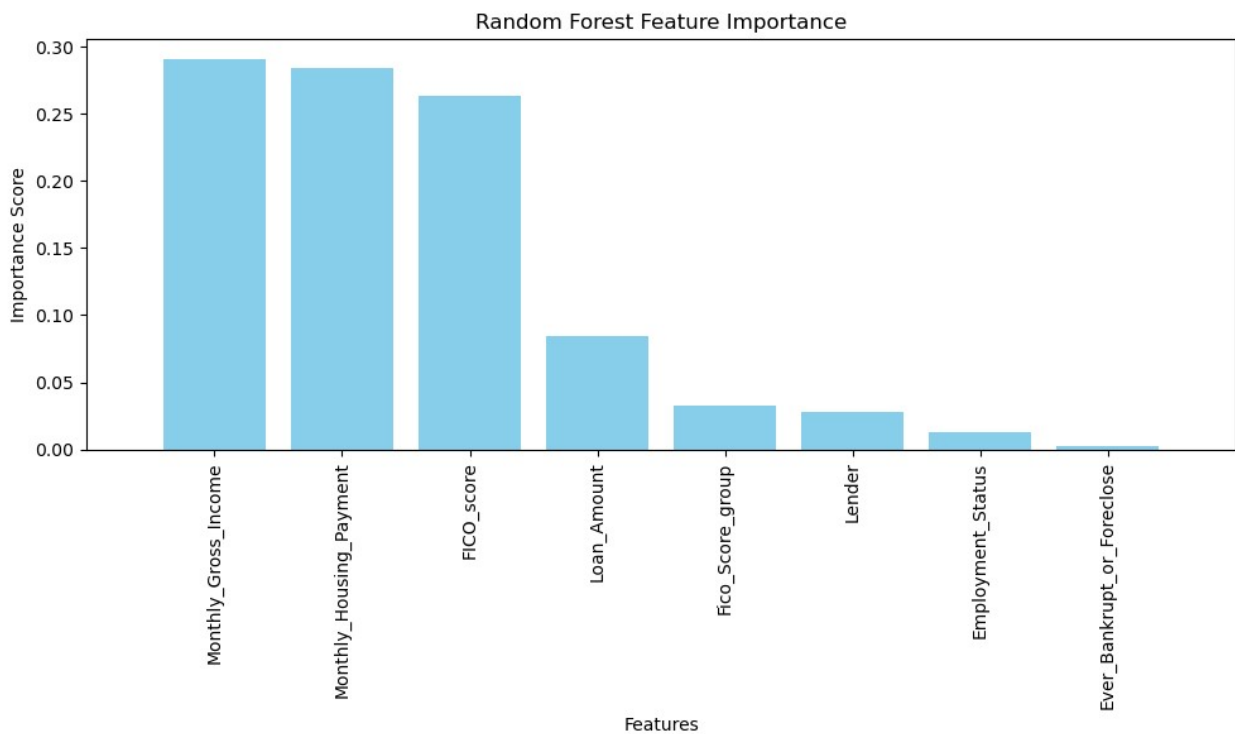
# Initialize and train the model
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Get feature importances
importances = rf.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# Visualize feature importances
```

```
plt.figure(figsize=(10, 6))
plt.title("Random Forest Feature Importance")
plt.bar(range(X.shape[1]), importances[indices], align="center", color = "skyblue")
plt.xticks(range(X.shape[1]), X.columns[indices], rotation=90)
plt.xlabel('Features')
plt.ylabel('Importance Score')
plt.xlim([-1, X.shape[1]])
plt.tight_layout()
plt.show()
```



## Exploring Lender Approval Rate

### Average approval rate

```
# Group the data by lender and calculate the average approval rate
lender_approval_rates = data_full.groupby('Lender')['Approved'].mean()

# Display the average approval rates for each lender
print("Average Approval Rates by Lender:")
print(lender_approval_rates)

# Visualize approval rates using a bar plot
plt.figure(figsize=(8, 6))
plt.bar(lender_approval_rates.index, lender_approval_rates.values,
color=['blue', 'green', 'orange'])
```

```
plt.xlabel('Lender')
plt.ylabel('Approval Rate')
plt.title('Approval Rates by Lender')
plt.ylim(0, 1) # Set y-axis limits to ensure proper visualization of approval rates
plt.grid(axis='y', linestyle='--', alpha=0.7) # Add gridlines for better readability
plt.show()
```

Average Approval Rates by Lender:

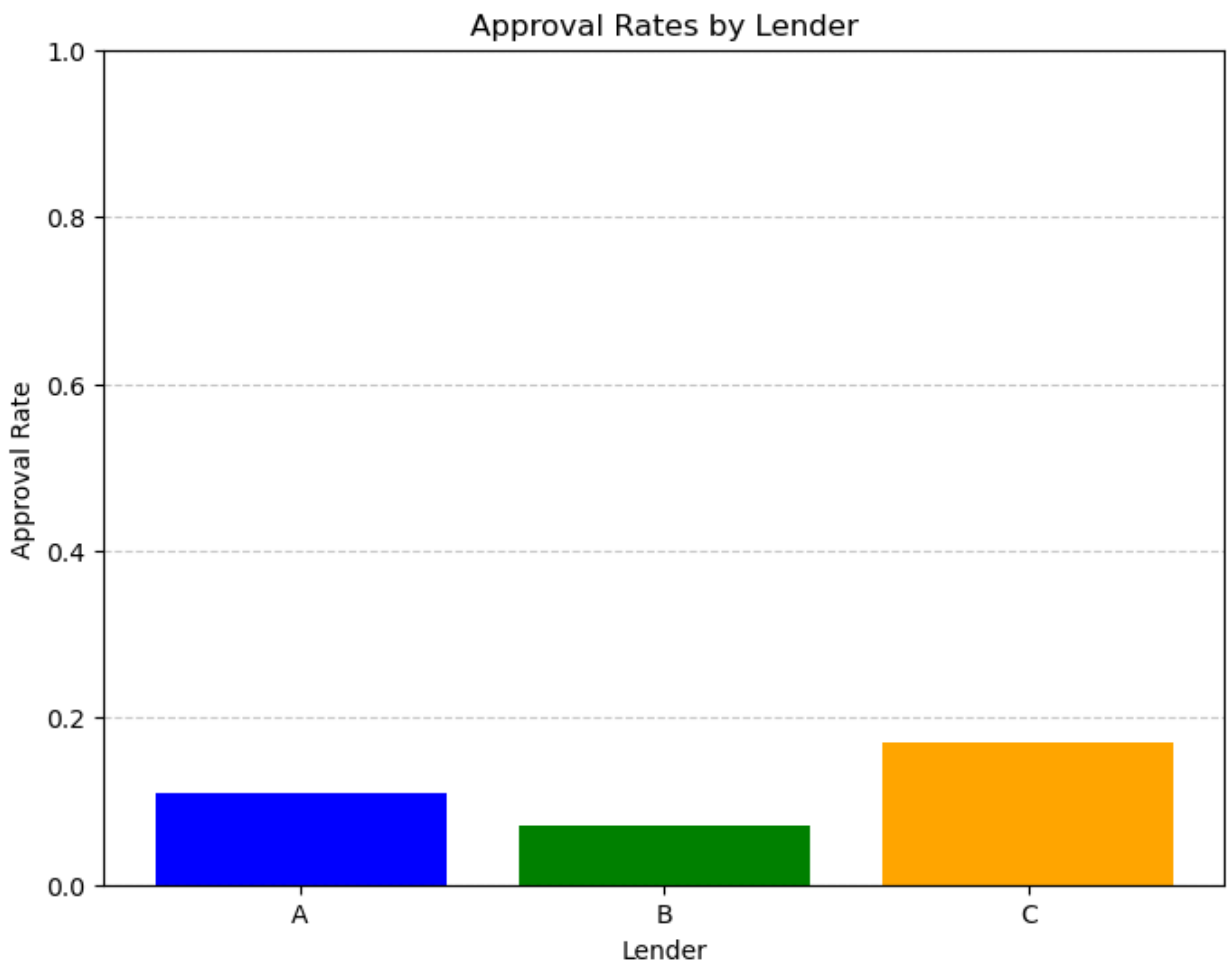
Lender

A 0.109655

B 0.071273

C 0.170571

Name: Approved, dtype: float64



Explore differences in costumers based on lender

*#Define relevant variables for comparison*

```
variables_of_interest = ['FICO_score', 'Monthly_Gross_Income',
```

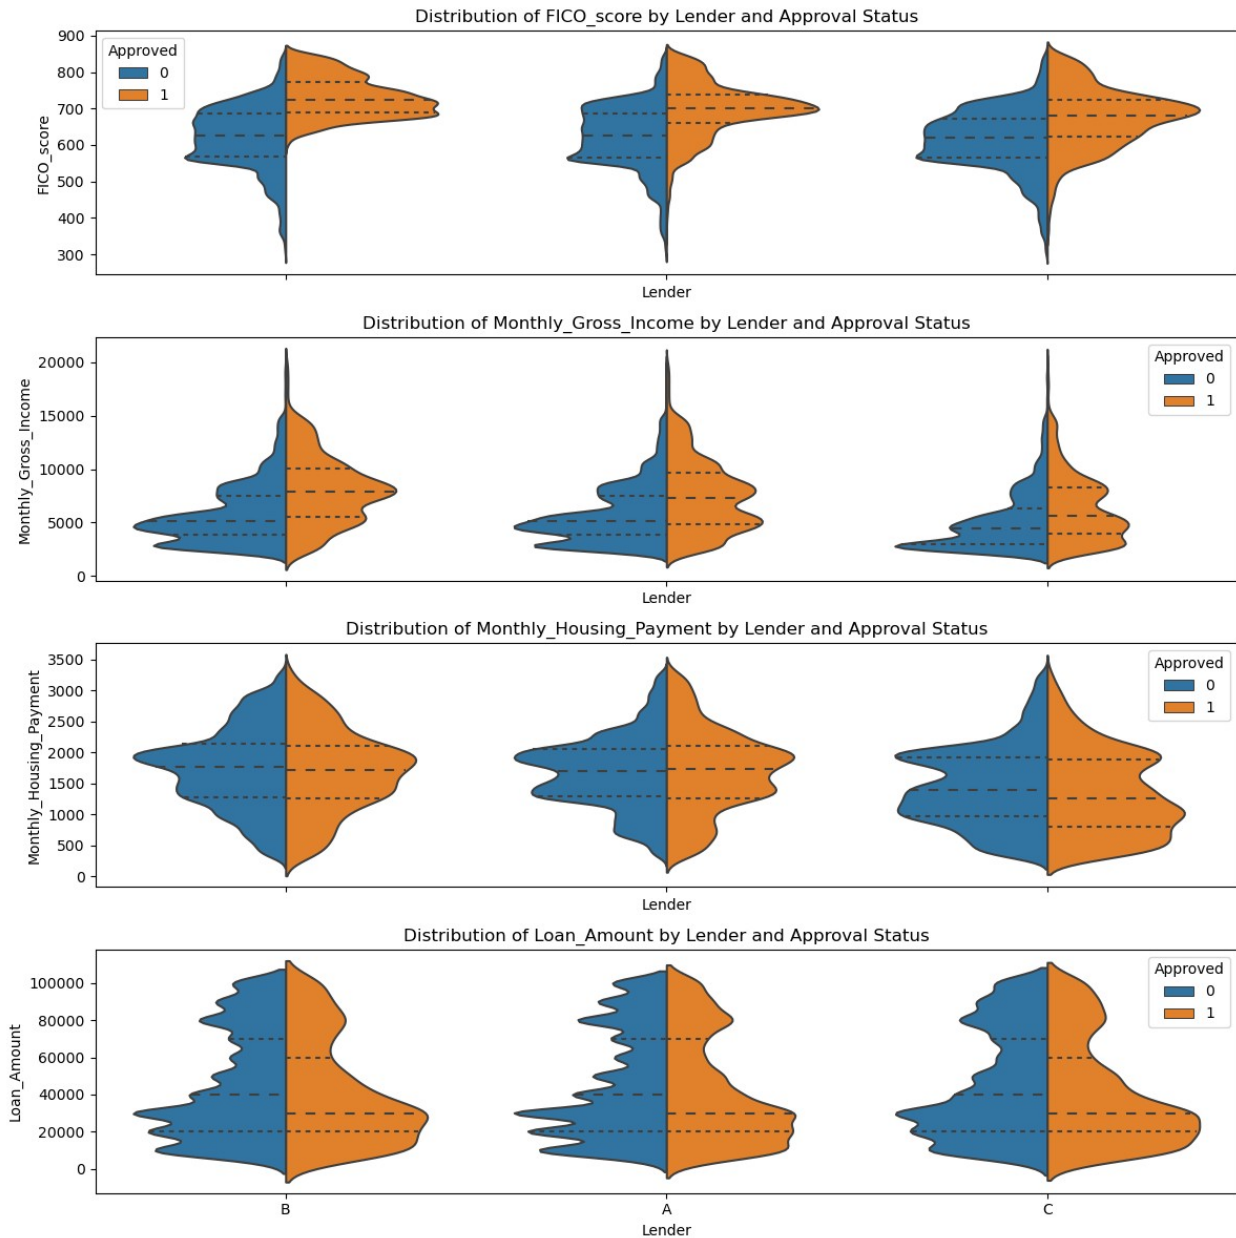
```
'Monthly_Housing_Payment','Loan_Amount']

# Set up the figure with subplots
fig, axes = plt.subplots(len(variables_of_interest), 1, figsize=(12,
12), sharex=True)

# Plot violinplots for each variable of interest
for i, variable in enumerate(variables_of_interest):
    sns.violinplot(data=data_full, x='Lender', y=variable,
hue='Approved', ax=axes[i], split=True, inner='quartile')
    axes[i].set_title(f'Distribution of {variable} by Lender and
Approval Status')
    axes[i].set_xlabel('Lender')
    axes[i].set_ylabel(variable)
    axes[i].legend(title='Approved')

# Adjust layout
plt.tight_layout()
plt.show()
```





## explore feature importance based on lender

```
# For demonstration, let's use logistic regression to identify
important variables
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

# Define features and target variable
features = ['FICO Score', 'Monthly Gross Income', 'Employment Status']
target = 'Approved'
```

```

# Fit a logistic regression model for each lender
lenders = data_full['Lender'].unique()

for lender in lenders:
    print(f"\nLogistic Regression Model for Lender: {lender}")

    # Filter data for the current lender
    X_train_lender = X_train[data_full['Lender'] == lender]
    y_train_lender = y_train[data_full['Lender'] == lender]
    X_test_lender = X_test[data_full['Lender'] == lender]
    y_test_lender = y_test[data_full['Lender'] == lender]

    # Create and fit logistic regression model
    model = LogisticRegression()
    model.fit(X_train_lender, y_train_lender)

    # Evaluate model performance
    train_accuracy = model.score(X_train_lender, y_train_lender)
    test_accuracy = model.score(X_test_lender, y_test_lender)
    print(f"Train Accuracy: {train_accuracy:.2f}")
    print(f"Test Accuracy: {test_accuracy:.2f}")

    coefficients = np.abs(model.coef_[0])
    feature_names = X.columns

    # Plotting
    plt.figure(figsize=(10, 6))
    plt.barh(feature_names, coefficients, color='skyblue')
    plt.xlabel('Coefficient Magnitude')
    plt.ylabel('Feature')
    plt.xticks(rotation = 90)
    plt.title('Feature Importances for lender '+lender)
    plt.show()

    # Generate classification report
    y_pred = model.predict(X_test_lender)
    print(classification_report(y_test_lender, y_pred))

```

```

Logistic Regression Model for Lender: B
Train Accuracy: 0.93
Test Accuracy: 0.93

```

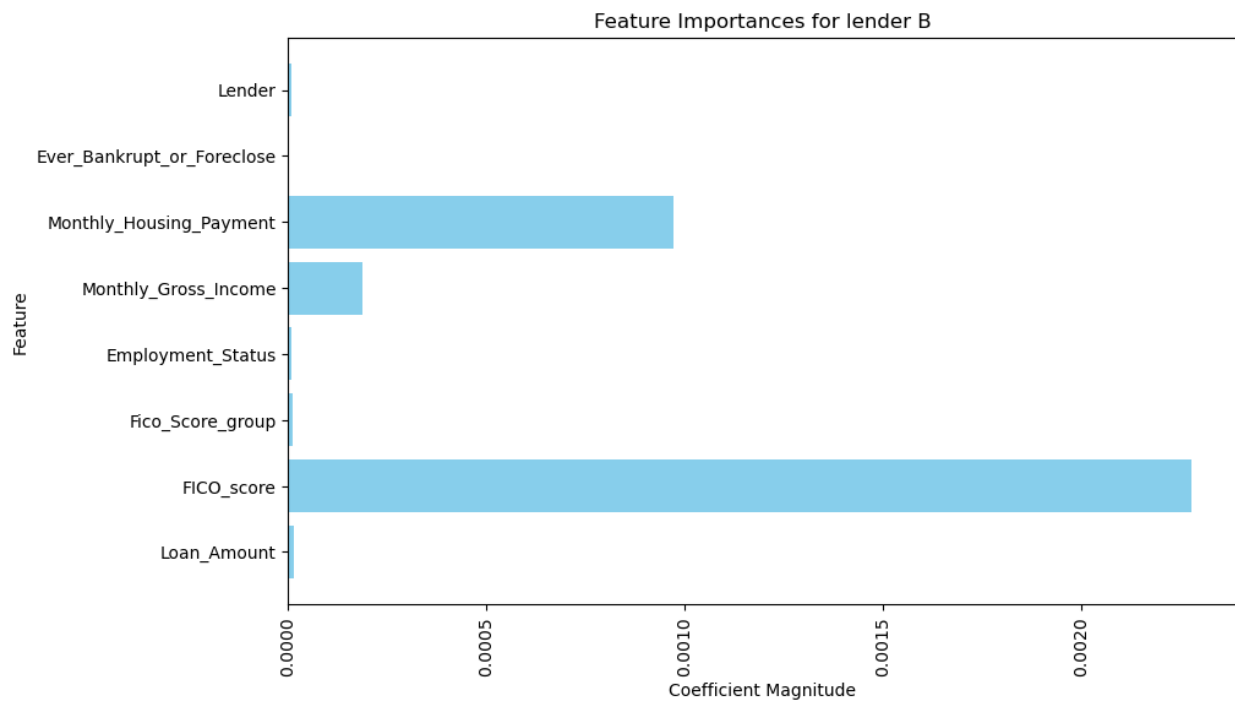
```

/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/
ipykernel_94964/3427714349.py:18: UserWarning: Boolean Series key will
be reindexed to match DataFrame index.
    X_train_lender = X_train[data_full['Lender'] == lender]
/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/ipykernel_94964/34277
14349.py:20: UserWarning: Boolean Series key will be reindexed to

```

```
match DataFrame index.
```

```
X_test_lender = X_test[data_full['Lender'] == lender]
```



	precision	recall	f1-score	support
0	0.93	1.00	0.96	7671
1	0.57	0.01	0.01	560
accuracy			0.93	8231
macro avg	0.75	0.50	0.49	8231
weighted avg	0.91	0.93	0.90	8231

Logistic Regression Model for Lender: A

Train Accuracy: 0.89

Test Accuracy: 0.89

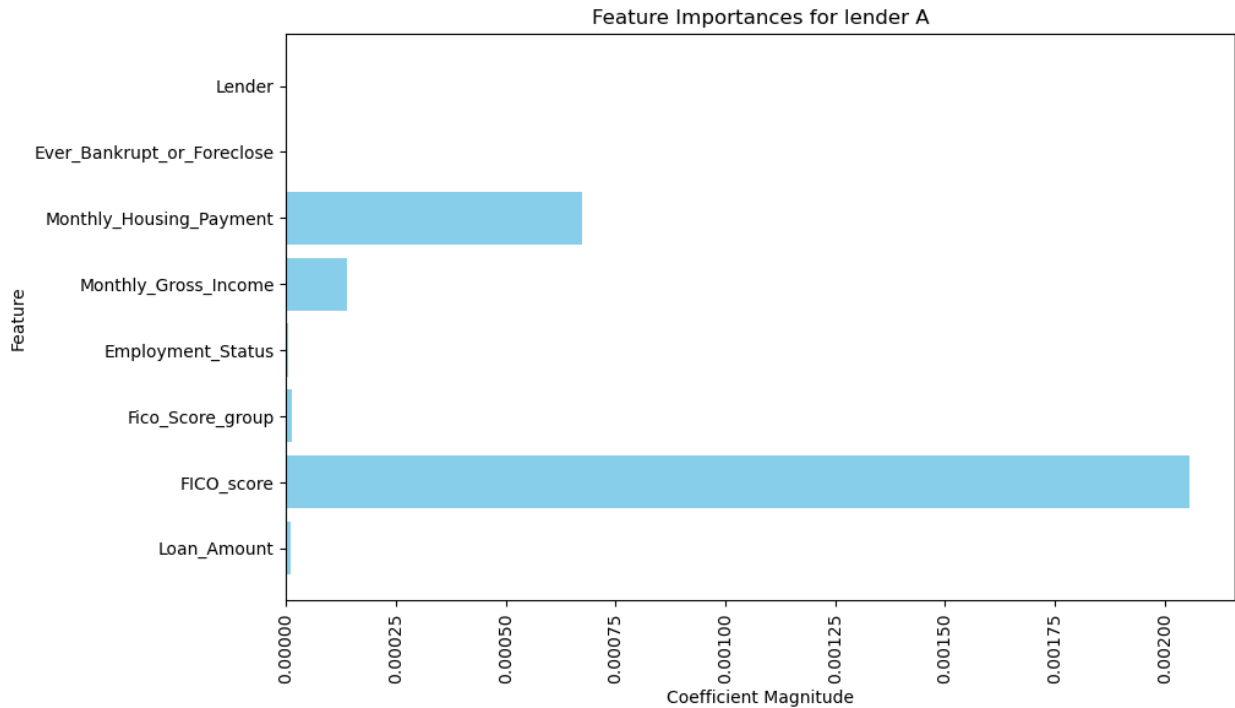
```
/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/
```

```
ipykernel_94964/3427714349.py:18: UserWarning: Boolean Series key will  
be reindexed to match DataFrame index.
```

```
X_train_lender = X_train[data_full['Lender'] == lender]
```

```
/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/ipykernel_94964/34277  
14349.py:20: UserWarning: Boolean Series key will be reindexed to  
match DataFrame index.
```

```
X_test_lender = X_test[data_full['Lender'] == lender]
```



	precision	recall	f1-score	support
0	0.89	1.00	0.94	14681
1	0.40	0.00	0.00	1780
accuracy			0.89	16461
macro avg	0.65	0.50	0.47	16461
weighted avg	0.84	0.89	0.84	16461

Logistic Regression Model for Lender: C

Train Accuracy: 0.83

Test Accuracy: 0.83


/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/

ipykernel\_94964/3427714349.py:18: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
X_train_lender = X_train[data_full['Lender'] == lender]
```

/var/folders/yn/ftl55p110bl7f4vqh175t6mh0000gn/T/ipykernel\_94964/3427714349.py:20: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
X_test_lender = X_test[data_full['Lender'] == lender]
```



	precision	recall	f1-score	support
0	0.83	1.00	0.90	4367
1	0.67	0.03	0.05	941
accuracy			0.83	5308
macro avg	0.75	0.51	0.48	5308
weighted avg	0.80	0.83	0.75	5308

## PART A

```
lenders = data_full['Lender'].unique()

imp_features = ['FICO_score', 'Loan_Amount', 'Monthly_Gross_Income',
'Monthly_Housing_Payment', 'bounty']

for feature in imp_features:
    group = data_full.groupby('Lender')[feature]
    mean = group.mean()
    print(mean)
```

Lender

A 630.125727

B 630.246509

C 625.500971

Name: FICO\_score, dtype: float64

Lender

```

A    45257.909091
B    45223.090909
C    45178.000000
Name: Loan_Amount, dtype: float64
Lender
A    5989.508036
B    5994.417745
C    5309.743143
Name: Monthly_Gross_Income, dtype: float64
Lender
A    1679.922945
B    1730.975636
C    1426.960286
Name: Monthly_Housing_Payment, dtype: float64
Lender
A    27.413636
B    24.945455
C    25.585714
Name: bounty, dtype: float64

```

## PART B

```

# Create a function to determine the best lender for each customer based on characteristics
def determine_best_lender(row):
    # Define criteria for selecting the best lender (e.g., based on FICO score, income, etc.)
    # For demonstration purposes, let's assume the lender with the highest average bounty is chosen
    best_lender = data.groupby('Lender')['bounty'].mean().idxmax()
    return best_lender

# Apply the function to determine the best lender for each customer
data['Best_Lender'] = data.apply(determine_best_lender, axis=1)

# Aggregate the total revenue per lender before matching
total_revenue_per_lender_before_matching = data.groupby('Lender')['bounty'].sum()

# Aggregate the total revenue per lender after matching
total_revenue_per_lender_after_matching = data.groupby('Best_Lender')['bounty'].sum()

# Calculate incremental revenue by subtracting total revenue per lender before matching from after matching
incremental_revenue = total_revenue_per_lender_after_matching - total_revenue_per_lender_before_matching

# Display the incremental revenue
#print("Incremental Revenue by Matching Customers to the Best

```

```
Lender:")  
print(incremental_revenue)
```

```
Best_Lender
```

```
2      0
```

```
Name: bounty, dtype: int64
```